

# Restoring Ankle Push-off Dynamics in Markerless Gait Analysis by Calibrated Deconvolution

F. Fer<sup>1</sup>

<sup>1</sup>Institut de Myologie — `myogait` technical report (proto-article)

August 2026


## Abstract

Markerless two-dimensional pose estimation reproduces hip and knee sagittal kinematics well but systematically under-reads the ankle range of motion (ROM), chiefly by flattening the fast plantar-flexion of push-off. We show that, on a joint-angle waveform, the pose estimator behaves as a subject-independent low-pass filter: it retains about 80% of the signal at 1 Hz but only  $\sim 20\%$  at 6–7 Hz. Modelling the estimator as a linear time-invariant system with transfer function  $H(f)$ , we estimate  $H$  once against synchronous optical motion capture (Vicon) on the BioCV dataset and invert it by Wiener deconvolution. The correction is applied as *mean restoration*: because deconvolution is linear, the systematic deficit lives in the mean cycle whereas the cycle-to-cycle spread is the clinical variability signal that must be preserved; we therefore deconvolve only the mean and add the resulting per-phase correction to every cycle. In a leave-one-subject-out validation over 9 subjects and 85 walking trials, the ankle ROM bias is halved ( $-10.6^\circ \rightarrow -5.3^\circ$ ), the absolute ROM error drops from  $10.8^\circ$  to  $6.7^\circ$ , and inter-cycle variability is preserved exactly ( $1.73^\circ$ ). The method requires no motion capture at run time, makes no healthy-gait assumption, and improves every subject.

## 1 Introduction

Optical marker-based motion capture is the clinical reference for joint kinematics but is costly, lab-bound and slow to set up. Video-based *markerless* pipelines built on deep pose estimators—OpenPose [1], and more recently high-resolution whole-body models such as Sapiens [2]—promise the same measurement from a single camera [3]. Concurrent-validity studies now consistently report that hip and knee sagittal angles reproduce the optical reference to a few degrees, while the *ankle* is the weakest joint [4, 5]. The dominant error is a flattened push-off: the real foot plantar-flexes sharply to propel the body, but the tracked foot landmarks do not follow that fast motion (Figure 1).

Two families of remedies exist. Learned *marker augmenters* map sparse video keypoints to dense anatomical markers using a network trained on motion capture, as in OpenCap [6]; parametric body-and-foot mesh models regularise the pose with a learned shape prior. Both are powerful but opaque and data-hungry. Here we ask whether a simple, interpretable, *calibrated* signal-processing correction is enough. Our contributions are: (i) the observation that the estimator acts as a subject-independent low-pass filter on joint-angle waveforms; (ii) a cadence-adaptive Wiener deconvolution that inverts it from a filter calibrated once against Vicon; and (iii) a *mean-restoration* scheme that recovers the systematic amplitude while leaving inter-cycle variability—itsself a clinical marker [7]—mathematically unchanged. We also report, for completeness, a set of geometric multi-view reconstructions that were tried and did not solve the ankle (Appendix A).



**Figure 1:** The deficit (subject P06). During push-off (shaded,  $\sim 55\text{--}72\%$  of the cycle) the optical reference reaches deep plantar-flexion, whereas the single lateral camera barely dips. The error is fast and phase-locked.

## 2 Method

### 2.1 Data and angle extraction

Video is processed by the `myogait` pipeline: per-axis normalisation, a 4 Hz Butterworth filter, 2-D sagittal joint angles, coordinate-based (Zeni) event detection, and time normalisation of each gait cycle to 101 points. The optical reference is the Visual3D-processed Vicon of the same trial, from which we recompute the ankle from the full 3-D markers. All angles are compared after removing the whole-cycle mean, so a definition offset between the two systems does not bias the comparison. Walking direction is resolved from several pan-invariant intra-frame body-orientation cues (foot-, nose- and toe-relative offsets) that remain valid when the camera pans to follow the subject, unlike the legacy hip-displacement heuristic. The final angle sign is then fixed by a flexion-positive canonicalisation that is itself independent of the detected direction; consequently the ankle metric reported here is invariant to the direction stage. We verified this directly: mirroring every landmark  $x$ -coordinate—the geometric equivalent of reversing the walking direction—reproduces the canonicalised ankle waveform to within  $1^\circ$ . Every number below is therefore identical whether the legacy or the newer detector is used.


### 2.2 The pose estimator as a linear filter

Write  $c(\phi)$  for the true (Vicon) ankle angle over the normalised cycle  $\phi$  and  $v(\phi)$  for the markerless one. Taking the discrete Fourier transform over the cycle, the ratio between the video and reference spectra at each stride harmonic  $k$  is, across subjects, a smooth and reproducible function of frequency that decays with  $k$  (Figure 2). We therefore model the estimator as a linear time-invariant system,

$$V(f) = H(f) C(f), \quad (1)$$

where  $H(f)$  is the complex transfer function of the tracker. Because  $H$  is a property of the estimator, not of the subject, it is estimated once and re-used. We sample it at the eight lowest harmonics with a Wiener (minimum mean-square) estimator over the training subjects,

$$\hat{H}_k = \frac{\langle V_k \overline{C_k} \rangle}{\langle |C_k|^2 \rangle}. \quad (2)$$



**Figure 2:** Estimated markerless transfer function  $|H(f)|$ : a low-pass. The tracker keeps  $\sim 80\%$  of the ankle signal near 1 Hz but only  $\sim 20\%$  by 6–7 Hz, exactly where the sharp push-off lives.

### 2.3 Cadence-adaptive inversion

At run time, for a trial of stride period  $T$ , harmonic  $k$  sits at frequency  $k/T$ ; we interpolate  $H$  there, so the correction adapts to each subject’s cadence and makes *no* assumption about a healthy gait pattern. The waveform is restored by regularised (Wiener) inverse filtering,

$$\hat{C}_k = V_k \frac{\overline{H_k}}{|H_k|^2 + \lambda}, \quad \lambda = 0.08, \quad (3)$$

the regulariser  $\lambda$  trading restoration against noise amplification.

### 2.4 Mean restoration

Deconvolution amplifies high-frequency noise; applied to each raw cycle it would inflate the very inter-cycle variability we wish to measure. Because the deconvolution operator  $\mathcal{D}$  is linear,  $\text{mean}(\mathcal{D}\{c_i\}) = \mathcal{D}\{\bar{c}\}$ : the systematic deficit lives in the *mean* cycle  $\bar{c}$ , while the deviations  $c_i - \bar{c}$  carry the variability. We therefore deconvolve only the mean, form the per-phase correction, and add it to every cycle:

$$\Delta(\phi) = \mathcal{D}\{\bar{c}\}(\phi) - \bar{c}(\phi), \quad c_i^{\text{corr}}(\phi) = c_i(\phi) + \Delta(\phi). \quad (4)$$

Every cycle receives the *same* correction  $\Delta$ , so the mean is restored to  $\mathcal{D}\{\bar{c}\}$  while the cycle-to-cycle spread is left mathematically unchanged. This is the property that makes the method safe for variability-based clinical readings.

### 2.5 Geometric reconstructions considered

Before adopting deconvolution we tested whether the ankle deficit was *geometric*—recoverable with additional calibrated cameras—rather than a limitation of the landmarks themselves. Two ingredients were combined.

**Skeleton locking (rigid-body constraints).** Human segments are rigid. We first estimate each bone length as the temporal median  $L_{ab} = \text{median}_t \|X_a(t) - X_b(t)\|$  and rebuild the limb by chaining unit segment directions scaled by  $L_{ab}$  from a triangulated pelvis root, so no frame can stretch or collapse a bone. The foot is treated as one rigid triangle (ankle–heel–toe): at each frame we fit the canonical foot template  $\{p_j\}$  to the observed points  $\{q_j\}$  by a rigid transform (a Procrustes / Kabsch solution),

$$(R^*, t^*) = \arg \min_{R \in SO(3), t} \sum_j \|Rp_j + t - q_j\|^2, \quad (5)$$

**Table 1:** Ankle metrics, leave-one-subject-out (9 subjects, 85 trials). Degrees.

| Metric (ankle)                | Before | <b>After</b> |
|-------------------------------|--------|--------------|
| ROM error  vs Vicon           | 10.76  | <b>6.71</b>  |
| ROM bias (video – Vicon)      | −10.61 | <b>−5.29</b> |
| Waveform RMSE vs Vicon        | 3.70   | <b>3.43</b>  |
| Inter-cycle SD (preserved)    | 1.73   | <b>1.73</b>  |
| Inter-patient SD (Vicon 5.18) | 3.63   | 4.75         |

and read the foot orientation off  $R^*$ . This removes the non-rigid jitter of the foot landmarks (their inter-marker distance was seen to vary by 20–26% in video versus 2–5% in Vicon), which cleans the ankle waveform (shape correlation with Vicon rose from 0.84 to 0.88) but does *not* restore the missing amplitude.

**Multi-view directional fusion.** A lateral camera constrains a segment’s antero-posterior and vertical components but is nearly blind to depth; a near-orthogonal (frontal) camera constrains the complementary axis. For a segment 3-D direction  $d$ , each camera  $c$  contributes one linear constraint: the observed 2-D direction  $u_c$  must be parallel to  $M_c d$ , with  $M_c$  the two in-plane rows of the camera rotation, i.e.  $(u_c^{\perp\top} M_c) d = 0$ . Two cameras give a  $2 \times 3$  homogeneous system whose null space is  $d$ ; chaining these directions with locked bone lengths yields a 3-D skeleton without triangulating the (noisy) foot position directly.


Both were evaluated but, as Appendix A shows, multi-view fusion *degraded* the ankle relative to a single lateral view—the frontal camera tracks the foot poorly—so the reported method uses one lateral camera plus the calibrated deconvolution.

### 3 Validation protocol

We use the BioCV dataset (University of Bath Research Data Archive [9]): synchronous Sapiens-2 markerless video (lateral camera) and Visual3D Vicon, 9 subjects, 85 walking trials. The transfer function is estimated **leave-one-subject-out**: for each test subject  $H$  is learned only on the others, so no test data enters the filter. Per trial and per side we report, for the ankle: the absolute ROM error and the signed bias versus Vicon; the waveform RMSE (mean-referenced); the inter-cycle SD (mean over phase of the standard deviation across a trial’s cycles); and, across subjects, the inter-patient SD of the mean ROM. Agreement is summarised by Bland–Altman bias and  $\pm 1.96$  SD limits of agreement (LoA).

### 4 Results

Figure 3 overlays, for every one of the nine subjects, the raw and corrected mean cycle (with the corrected inter-cycle band) and the Vicon mean. The corrected curve deepens push-off toward the reference in every subject, and—crucially—the corrected band has the same width as the raw band. Table 1 gives the pooled metrics: the ROM bias is halved and the absolute error drops by 38%, while the inter-cycle SD is unchanged to two decimals. The Bland–Altman plots (Figure 4) show the bias moving toward zero and the lower limit of agreement tightening. The correction improves *every* subject (Figure 5), and the between-subject spread of ankle ROM moves *toward* the Vicon value ( $3.6^\circ \rightarrow 4.7^\circ$ , Vicon  $5.2^\circ$ ): real inter-patient differences are recovered, not smoothed away.




**Figure 3:** All nine subjects (per panel: raw ROM  $\rightarrow$  restored / Vicon). Raw video (grey, dashed), restored mean  $\pm$  inter-cycle SD (orange), Vicon (teal). The restored curve deepens push-off toward Vicon in every subject while the band width is preserved.


## 5 Discussion

The correction halves a systematic, phase-locked bias with a single lateral camera, one calibrated filter and about ten lines of code at run time. Three properties make it suitable for clinical, and specifically pathological, use. *Variability is preserved* by construction, so gait variability is neither smoothed nor inflated. *Inter-patient differences are restored, not erased*, which matters when the between-patient signal is itself the finding. And the filter is applied in Hz at the subject’s own cadence, so there is *no healthy-gait prior* that could hallucinate a normal pattern onto a pathological one. It generalises: leave-one-subject-out, all nine subjects improved.

**Limitations.** (i) The residual bias ( $\sim 5^\circ$ ) is information the estimator never captured; a linear inverse filter can amplify what remains but cannot recreate what is absent—closing the gap further would require a learned non-linear augmenter [6]. (ii) The correction restores *amplitude* (ROM); the point-wise RMSE improves only modestly, as some shape error persists. (iii)  $H$  was estimated on one tracker and one dataset of healthy adults; a different tracker or population should re-estimate  $H$  (the calibration, not the method, is specific). (iv) Multi-view fusion did not help the ankle (Appendix A), so a single lateral view plus this correction is the recommended configuration. A definitive clinical



**Figure 4:** Bland–Altman on ankle ROM. The bias (solid) moves toward zero and the lower limit of agreement (dashed) tightens after restoration.



**Figure 5:** Per-subject absolute ankle ROM error, before and after. Every subject improves.

claim will require concurrent validation on a pathological cohort.


## 6 Conclusion

Treating a markerless pose estimator as a calibrated linear system and inverting it—while restoring only the systematic mean—recovers a clinically meaningful part of the ankle push-off that 2-D video loses, without motion capture at run time and without touching gait variability. The method ships in `myogait` as the opt-in `analyze_gait(restore_ankle_dynamics=True)`.

## A Alternative reconstructions and ablation

On the one subject with four calibrated views plus Vicon, we scored ten multi-view reconstruction strategies by mean absolute ROM error against Vicon (Figure 6). Weighted or averaged fusion never beat the best single view; the strongest variant fused orthogonal image axes with locked bone lengths (Section 2.5), yet even it plateaued, because the push-off is largely absent from the foot landmarks—no geometric method fully recovers it.

The step-by-step ablation on the ankle (Table 2) shows that deconvolution on a single lateral camera is the best single step; multi-camera fusion degrades the ankle and the methods do not stack. The




**Figure 6:** Ten multi-view reconstruction strategies (subject P06), mean |ROM error| vs Vicon. Teal: clear improvement (geometry-limited plateau); grey: failed.

raw per-cycle deconvolution used in this ablation inflates the inter-cycle SD; the mean-restoration variant of the main text removes that side-effect while keeping the accuracy gain.

**Table 2:** Ablation on the ankle (subject P06; Vicon ROM  $\approx 38^\circ$ ). Degrees.

| Configuration                      | ROM         | RMSE        | inter-cycle SD |
|------------------------------------|-------------|-------------|----------------|
| Single lateral camera              | 24.7        | 4.40        | 4.71           |
| + <b>deconvolution (this work)</b> | <b>32.7</b> | <b>3.37</b> | <b>5.84</b>    |
| Multi-camera reconstruction        | 20.8        | 5.80        | 5.70           |
| + skeleton/foot locking            | 20.0        | 5.77        | 5.34           |
| + locking + deconvolution          | 23.9        | 6.35        | 6.55           |

Figure 7 shows the correction *effect* of each alternative on the mean ankle waveform: only the deconvolution (orange) deepens the push-off toward Vicon, whereas the multi-camera reconstruction and skeleton/foot locking leave it as flat as—or flatter than—the single camera. This is the visual counterpart of the table: geometry does not restore the missing amplitude, the calibrated filter does.



**Figure 7:** Effect of each alternative on the mean ankle waveform (subject P06, foot-flat referenced). Only the deconvolution (this work) restores the push-off toward Vicon; multi-camera reconstruction and skeleton/foot locking do not.

## References

- [1] Z. Cao, G. Hidalgo, T. Simon, S.-E. Wei, Y. Sheikh. OpenPose: Realtime multi-person 2D pose estimation using part affinity fields. *IEEE TPAMI* 43(1):172–186, 2021.
- [2] R. Khrodgar, T. Bagautdinov, J. Martinez, et al. Sapiens: Foundation for human vision models. *ECCV*, 2024.
- [3] S. L. Colyer, M. Evans, D. P. Cosker, A. I. T. Salo. A review of the evolution of vision-based motion analysis and the integration of advanced computer vision methods towards developing a markerless system. *Sports Medicine – Open* 4:24, 2018.
- [4] R. M. Kanko, E. Laende, et al. Concurrent assessment of gait kinematics using marker-based and markerless motion capture. *Journal of Biomechanics* 127:110665, 2021.
- [5] L. Needham, M. Evans, D. P. Cosker, et al. The accuracy of several pose estimation methods for 3D joint centre localisation. *Scientific Reports* 11:20673, 2021.
- [6] S. D. Uhlich, A. Falisse, Ł. Kidziński, et al. OpenCap: Human movement dynamics from smartphone videos. *PLoS Computational Biology* 19(10):e1011462, 2023.
- [7] J. M. Hausdorff, D. A. Rios, H. K. Edelberg. Gait variability and fall risk in community-living older adults. *Archives of Physical Medicine and Rehabilitation* 82(8):1050–1056, 2001.
- [8] D. A. Winter. *Biomechanics and Motor Control of Human Movement*, 4th ed. Wiley, 2009.
- [9] M. Evans, L. Needham, S. Colyer, et al. BioCV dataset. University of Bath Research Data Archive, doi:10.15125/BATH-01258.